

# STAT 716 - Class 9/10 - 2018-11-05

*Vitaly Druker*

## Non Linearity

### Polynomial regression

global function - this can be an issue - especially at the outer edges

### Step Functions

Step Functions localize everything, cut points categorical dummy variables

### Basis functions

general functions

## Regression Splines

piecewise polynomial regression

show page 272

piecewise cubic regression where to put knots?

spline - continuous on second derivative

Natural splines

the function is required to be linear at the boundary (in the region where  $X$  is smaller than the smallest knot, or larger than the largest knot).

## Generalized Additive Models

We have discussed regular linear regression. We can standardize into GAM with the following equation:

$$y_i = \sum_{j=1}^p f_j(x_{ij}) + \epsilon_i$$

## Lab

```
library(ISLR)
first_fit <- lm(wage ~ poly(age, 4), data = Wage)

coef(summary(first_fit))
```

```
##               Estimate Std. Error   t value    Pr(>|t|)
## (Intercept)   111.70361  0.7287409 153.283015 0.000000e+00
## poly(age, 4)1  447.06785 39.9147851  11.200558 1.484604e-28
## poly(age, 4)2 -478.31581 39.9147851 -11.983424 2.355831e-32
## poly(age, 4)3  125.52169 39.9147851   3.144742 1.678622e-03
## poly(age, 4)4 -77.91118 39.9147851  -1.951938 5.103865e-02
```

```
coef(summary(lm(wage ~age + I(age^2) + I(age^3) + I(age^4),
                data = Wage)))
```

```
##               Estimate Std. Error   t value    Pr(>|t|)
## (Intercept) -1.841542e+02 6.004038e+01 -3.067172 0.0021802539
## age          2.124552e+01 5.886748e+00  3.609042 0.0003123618
## I(age^2)     -5.638593e-01 2.061083e-01 -2.735743 0.0062606446
## I(age^3)      6.810688e-03 3.065931e-03  2.221409 0.0263977518
## I(age^4)     -3.203830e-05 1.641359e-05 -1.951938 0.0510386498
```

Why don't the match? See here for an example: <https://stackoverflow.com/questions/29999900/poly-in-lm-difference-between-raw-vs-orthogonal>

```
coef(summary(lm(wage ~poly(age, 4, raw = T),
                data = Wage)))
```

```
##               Estimate Std. Error   t value    Pr(>|t|)
## (Intercept)   -1.841542e+02 6.004038e+01 -3.067172 0.0021802539
## poly(age, 4, raw = T)1  2.124552e+01 5.886748e+00  3.609042 0.0003123618
## poly(age, 4, raw = T)2 -5.638593e-01 2.061083e-01 -2.735743 0.0062606446
## poly(age, 4, raw = T)3  6.810688e-03 3.065931e-03  2.221409 0.0263977518
## poly(age, 4, raw = T)4 -3.203830e-05 1.641359e-05 -1.951938 0.0510386498
```

## Cuts

```
coef(summary(lm(wage ~cut(age, 4),
                data = Wage)))
```

```
##               Estimate Std. Error   t value    Pr(>|t|)
## (Intercept)    94.158392  1.476069 63.789970 0.000000e+00
## cut(age, 4) (33.5,49] 24.053491  1.829431 13.148074 1.982315e-38
## cut(age, 4) (49,64.5] 23.664559  2.067958 11.443444 1.040750e-29
## cut(age, 4) (64.5,80.1] 7.640592  4.987424  1.531972 1.256350e-01
```

## Splines

```
library(splines)
library(dplyr)

# cubic splines

lm(wage~bs(age,knots=c(25,40,60)),data=Wage) %>%
  coef
```

```
##               (Intercept) bs(age, knots = c(25, 40, 60))1
##               60.49371                3.98050
## bs(age, knots = c(25, 40, 60))2 bs(age, knots = c(25, 40, 60))3
```

```
##                44.63098                62.83879
## bs(age, knots = c(25, 40, 60))4 bs(age, knots = c(25, 40, 60))5
##                55.99083                50.68810
## bs(age, knots = c(25, 40, 60))6
##                16.60614
```

We have 3 knots - how many degrees of freedom?

If it's piecewise cubic how many degrees of freedom?  $4 + 4 + 4 + 4 = 16$

What do splines add? 1. Continuity 2. Continuity on 1st derivative 3. Continuity on 2nd derivative

Each constraint is a degree of freedom

remove  $3 + 3 + 3 = 9$

Remaining = 7 which is what we see

```
lm(wage~bs(age, df = 6),data=Wage) %>%
  coef
```

```
##      (Intercept) bs(age, df = 6)1 bs(age, df = 6)2 bs(age, df = 6)3
##      56.31384      27.82400      54.06255      65.82839
## bs(age, df = 6)4 bs(age, df = 6)5 bs(age, df = 6)6
##      55.81273      72.13147      14.75088
```

```
attributes(bs(Wage$age, df = 6))
```

```
## $dim
## [1] 3000    6
##
## $dimnames
## $dimnames[[1]]
## NULL
##
## $dimnames[[2]]
## [1] "1" "2" "3" "4" "5" "6"
##
##
## $degree
## [1] 3
##
## $knots
##      25%    50%    75%
## 33.75 42.00 51.00
##
## $Boundary.knots
## [1] 18 80
##
## $intercept
## [1] FALSE
##
## $class
## [1] "bs"      "basis"   "matrix"
```

```
lm(wage~bs(age),data=Wage) %>%
  coef
```

```
## (Intercept)    bs(age)1    bs(age)2    bs(age)3
```

```
##      58.68868    102.64368    48.76204    40.80330
```

```
lm(wage~ns(age, knots=c(25,40,60)),data=Wage) %>% coef() %>% length()
```

```
## [1] 5
```

```
lm(wage~bs(age, knots=c(25,40,60)),data=Wage) %>% coef() %>% length()
```

```
## [1] 7
```

```
lm(wage~ns(age, knots=c(25,40,60)),data=Wage) %>% coef() %>% length()
```

```
## [1] 5
```

How many degrees of freedom do we get from add the natural spline constraint

only 2 degrees of freedom on the ends

- 4 degrees of freedom

$2 + 4 + 4 + 2$

- $2 - 3 - 2$

```
2 + 4 + 4 + 2 +  
- 2 - 3 - 2
```

```
## [1] 5
```

```
ns_mod <- lm(wage~ns(age, df = 6),data=Wage)
```

```
bs_mod <- lm(wage~bs(age, df = 6),data=Wage)
```

```
poly_mod <- lm(wage~poly(age,6),data=Wage)
```

```
length(coef(ns_mod))
```

```
## [1] 7
```

```
length(coef(bs_mod))
```

```
## [1] 7
```

```
length(coef(poly_mod))
```

```
## [1] 7
```

```
age_df <- data.frame(age = seq(min(Wage$age), max(Wage$age), length.out = 100))
```

```
predict_se_fit <- function(model, newdata,...){
```

```
  predictions <- predict(model, newdata, se.fit = T)
```

```
  data.frame(  
    fit = predictions$fit,  
    se.fit = predictions$se.fit  
  )  
}
```

```
all_models <- bind_rows(  
  age_df %>%  
    nest() %>%  
    mutate(mod = map(data, ~predict_se_fit(ns_mod, .)),  
           model = "ns") %>%
```

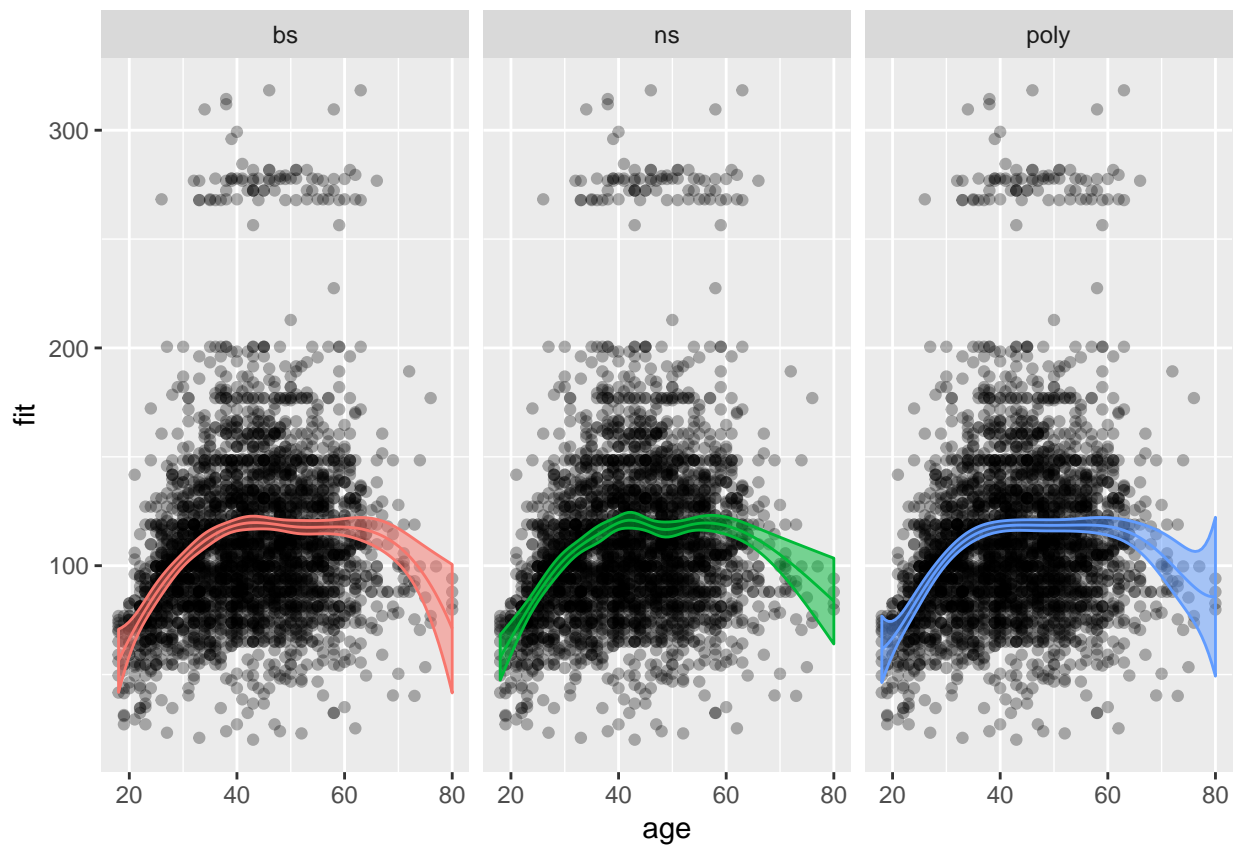
```

  unnest(),
age_df %>%
  nest() %>%
  mutate(mod = map(data, ~predict_se_fit(bs_mod, .)),
         model = "bs") %>%
  unnest(),
age_df %>%
  nest() %>%
  mutate(mod = map(data, ~predict_se_fit(poly_mod, .)),
         model = "poly") %>%
  unnest()

) %>%
  mutate(fit_low = fit - 2*se.fit,
         fit_high = fit + 2*se.fit)

all_models %>%
  ggplot(aes(x = age, y = fit, fill = model, color = model)) +
  geom_point(aes(x = age, y = wage), data = Wage, inherit.aes = F, alpha = 0.3) +
  geom_ribbon(aes(ymax = fit_high, ymin = fit_low), alpha = 0.5) +
  geom_line() +
  facet_wrap(~model) +
  theme(legend.position = "none")

```



## Smoothing Splines

$$\sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt$$

Loss + penalty function

2nd derivative is a measure of roughness

The function  $g(x)$  that minimizes (7.11) can be shown to have some special properties: it is a piecewise cubic polynomial with knots at the unique values of  $x_1, \dots, x_n$ , and continuous first and second derivatives at each knot. Furthermore, it is linear in the region outside of the extreme knots. In other words, the function  $g(x)$  that minimizes (7.11) is a natural cubic spline with knots at  $x_1, \dots, x_n$ !

effective parameters

$n$  to 2 as  $\lambda$  goes to infinity

you can use loocv to calc

$$RSS_{cv} = \sum_{i=1}^n \left( \frac{y_i - \hat{g}_\lambda(x_i)}{1 - \{S_\lambda\}_{ii}} \right)^2$$

$\hat{g}$  is the fitted values for the full spline

```
cv_smooth_spline <- with(Wage, smooth.spline(age, wage, cv=TRUE))
```

```
## Warning in smooth.spline(age, wage, cv = TRUE): cross-validation with non-  
## unique 'x' values seems doubtful
```

```
gcv_smooth_spline <- with(Wage, smooth.spline(age, wage))
```

```
library(gam)
```

```
## Loading required package: foreach
```

```
##
```

```
## Attaching package: 'foreach'
```

```
## The following objects are masked from 'package:purrr':
```

```
##
```



```
##      accumulate, when
```


```
## Loaded gam 1.16
```

```
gam_mod <- gam(wage ~ s(year) + s(age) + education, data = Wage)
```

```
gam_mod <- gam(wage ~ s(year, df = 5) + s(age) + education, data = Wage)
```

```
plot(gam_mod, se = T)
```






## Local Regression

weighted regression,


choose points within  $x_0$ . can be linear or w/e

```
data.frame(age = Wage$age,
span_2 = predict(loess(wage ~ age, span = .2, data = Wage))) %>%
  ggplot(aes(x = age, y = span_2)) +
  geom_point()
```





```
data.frame(age = Wage$age,  
span_5 = predict(loess(wage ~ age, span = .5, data = Wage))) %>%  
  ggplot(aes(x = age, y = span_5)) +  
  geom_point()
```



## Homework

Read Chapter 7 Question 9

## Missing Data

### Mean/Median Imputation

### Modeling

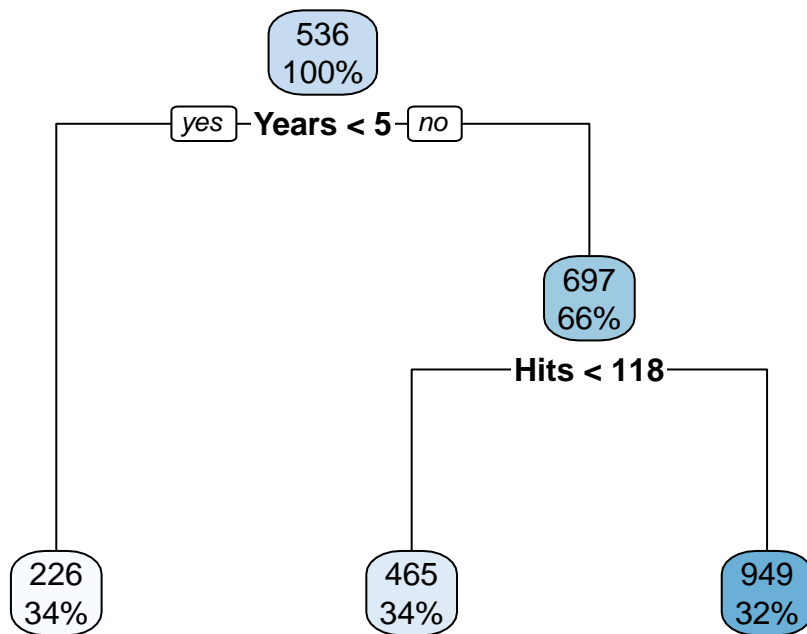
<https://topepo.github.io/caret/pre-processing.html#impute>

## Tree Based Methods


This is a method of splitting up the predictor space into sections.

```
library(rpart)
library(rpart.plot)

rpart(Salary ~ Years + Hits, Hitters, control = rpart.control(maxdepth = 2)) %>%
  rpart.plot()
```



Hitters %>%  
`ggplot(aes(x = Years, y = Hits)) + geom_point()`



Node - is a decision point

Terminal Node/Leaf - the final end of the tree or the predictor Internal Node

How does prediction work?

How do we make a tree?

1. Divide the predictor space into non overlapping regions. This means over *all* predictors.
2. Make a prediction for everyone that falls into a bucket

Ideally we could look at every single subset of features but that is not computationally feasible (akin to best subset) so we take a top down approach. Using recursive inarty splitting.

1. Pick a predictor and cut into 2 pieces by way of reducing RSS.

We do this by minimizing the joint RSS (it's the same as a step function). What does the RSS look like?

We then split one of those two regions to find the best next split.

This continues until a stoping criteria is reached (there are many different examples of stopping criteria that we will see in the lab)

Unfortunately this is a high variance method - we can't keep going or we will overfit the data.

## Cost Complexity Pruning

What can we penalize?

Prune back for a value that has a penalty of  $\alpha * |T|$  (the number of terminal nodes).

```
library(caret)
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
d_hit <- Hitters %>%
```


```
  select(Salary, Years, Hits) %>%
```

```
  filter(complete.cases(.))
```

```
train(Salary ~ Years + Hits, d_hit, method = "rpart", tuneLength = 10) %>% plot
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info =
```

```
## trainInfo, : There were missing values in resampled performance measures.
```



## Classification Trees

What error do we use?

Classification Error?

It is not sensitive enough Gini Coefficient can be used:

$$G = \sum_i^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$

Cross Entropy

$$\sum_i^K \hat{p}_{mk} \log(\hat{p}_{mk})$$

You can use either Gini or Cross Entropy to build trees - but use classification error to prune them

Benefits

Very easy to explain Nice to display Handle both qual and quant variable easily.

Issues

Hi variability - How does pruning help with variability of the tree? Where is the higher variability?

## Dealing with Issues of Trees

### Bagging

Bootstrap Aggregation (or bagging) can help methods that have a high variance without losing too much bias.

OOB Error Estimation - Natural cross validation.

Is the number of bootstraps a tuning parameter?

How do we interpret the models? You can look at the reduction of RSS when a variable is added and average it to see what's most important.

### Random Forests

Try to decorrelate the trees by using a random sample of  $m$  predictors

$$m = \sqrt{p}$$

If there is a strong predictor it will show up at the top of each bagged tree so we try not to use it every time.

### Homework

Read Chapter 8 Chapter 8 question 7